



Spec-o3 · Qwen3.5-4B Evaluation

Same 256 dev examples · Base / SFT (epoch 2) / RL (epoch 1) · shared answer extraction

(a) Classification Performance



(b) Format Compliance & Correct Answers



Generation caps: Base 6,132; SFT 2,048; RL 6,144 tokens per turn. Base alone uses an extra format reminder.

Panel (a) starts at zero; panel (b) uses separate y-axes and connects only the evaluated models.