


Spec-o3 · Qwen3.5-4B / 9B Evaluation


Same 256 dev examples · 4B Base / SFT epoch 2 / RL epochs 1 & 3 · 9B RL epoch 3

9B vs 4B at RL epoch 3: +16 correct · +6.25 pp accuracy · +5.01 pp Macro F1


(a) Accuracy




(b) Macro F1



(c) Strict Format Compliance



(d) Correct Answers



Generation caps: 4B Base 6,132; 4B SFT 2,048; all RL checkpoints 6,144 tokens per turn.

Base alone uses an extra format reminder. All runs use the same answer extraction and a 16,384-token context cap.